



Embedded Systems

Prof. Myung-Eui Lee (F-102)

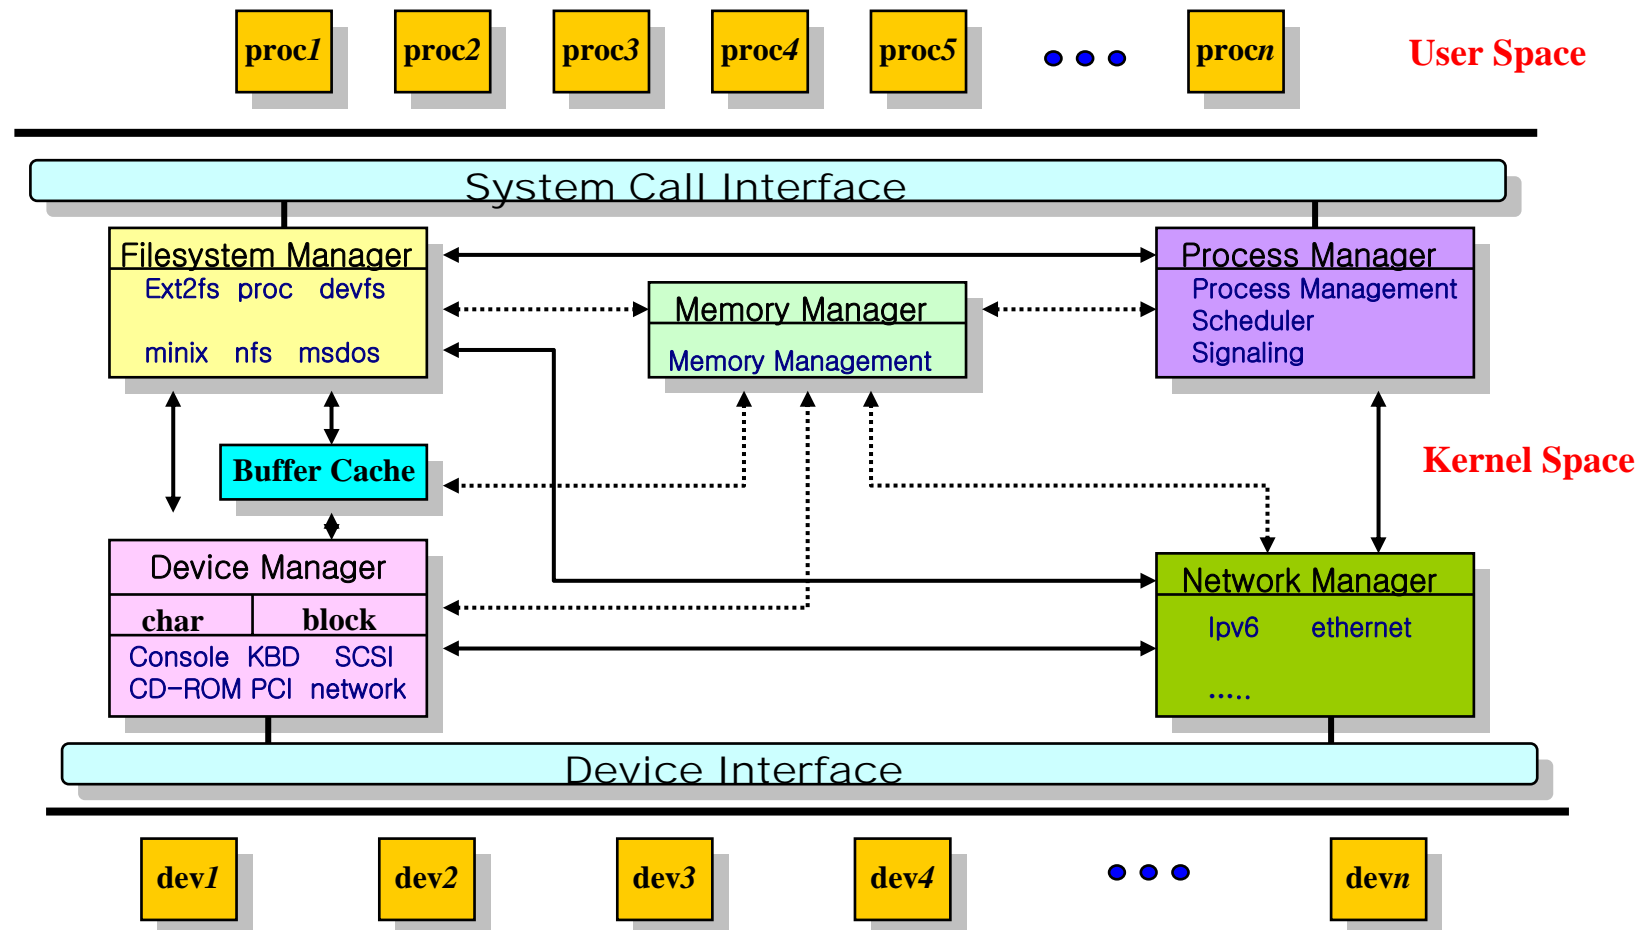
melee@kut.ac.kr





Linux Structure

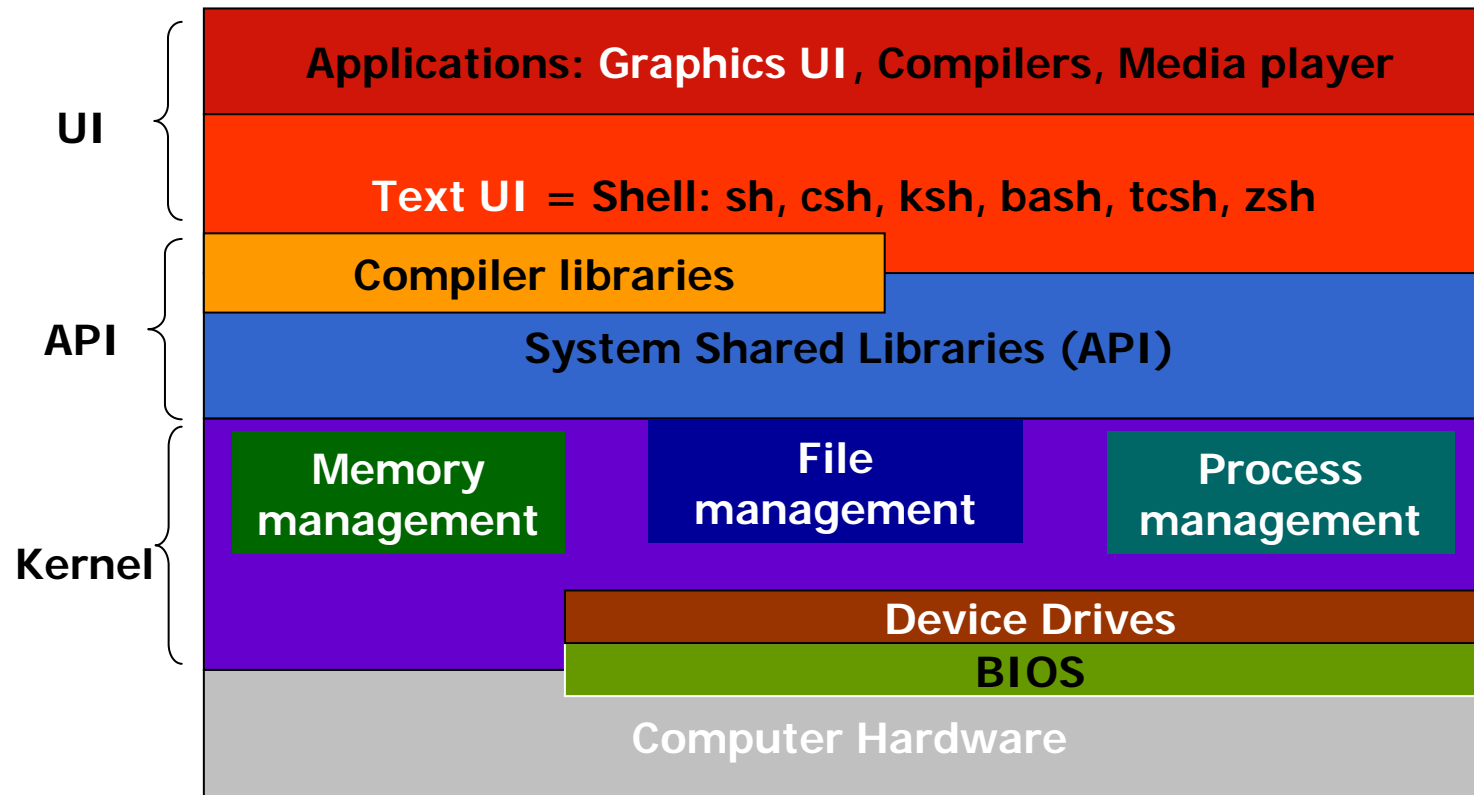
- Linux Structure





Linux Structure

- Linux Structure

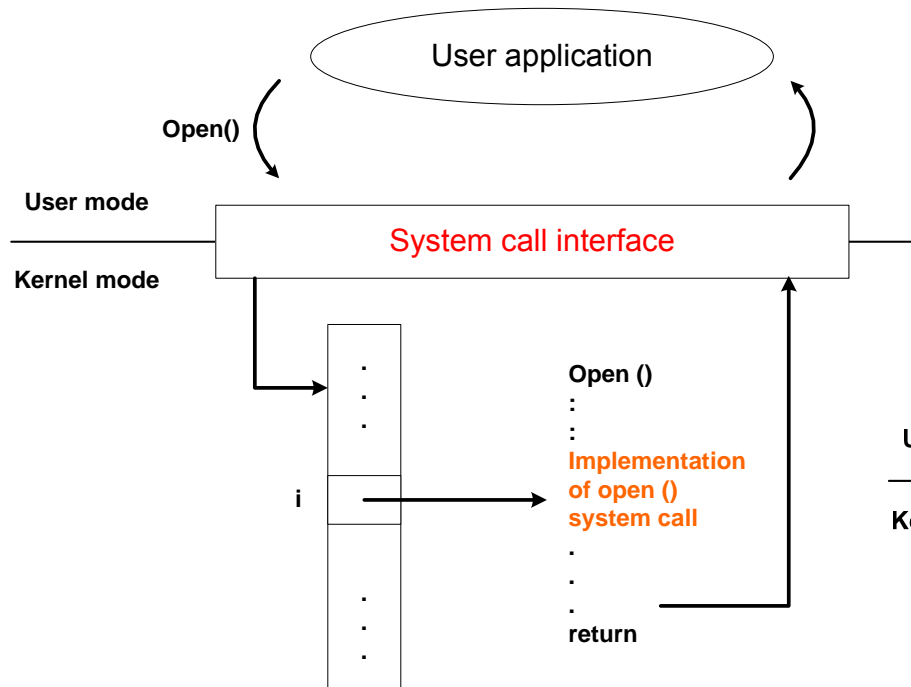




System Call Interface

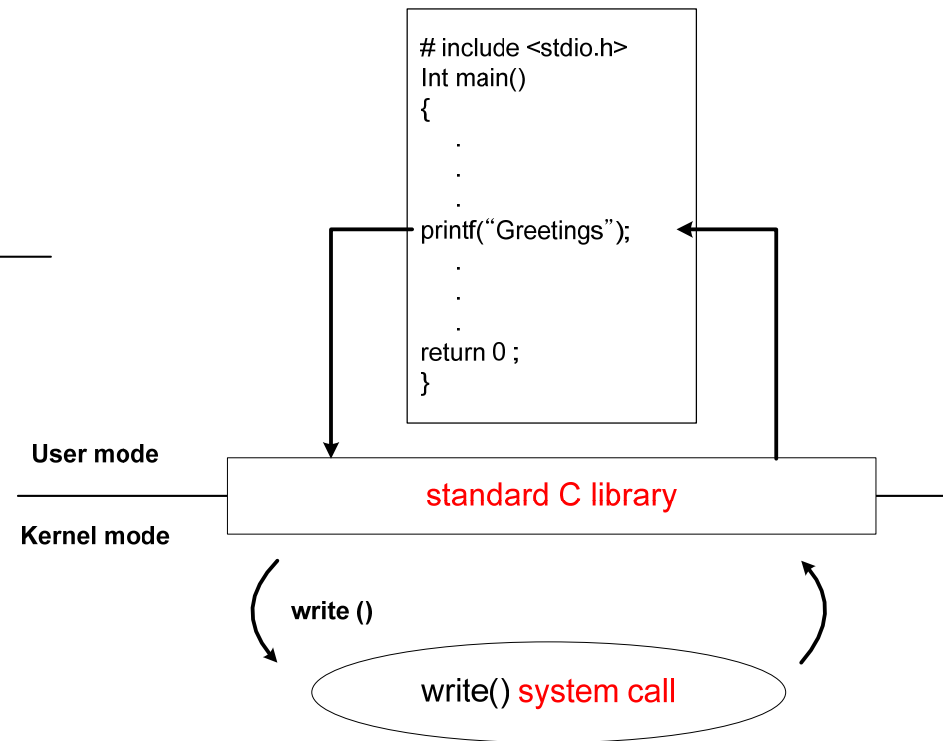
System Call

API – System call – OS Relationship



Standard C Lib.

C program invoking printf() library call, which calls write() system call





System Call

- **System calls define the programmer interface to Linux system**
 - » Interface between user-level processes and hardware devices.
 - ◆ CPU, memory, disks etc.
 - » **Make programming easier:**
 - ◆ Let kernel take care of hardware-specific issues.
 - » **Increase system security:**
 - ◆ Let kernel check requested service via system call.
 - » **Provide portability:**
 - ◆ Maintain interface but change functional implementation
- **Roughly five categories of system calls in Linux**
 - » Process control
 - » File management
 - » Device management
 - » Information maintenance
 - » Communications



System Call

- **Invoked by executing `int` or `swi` instruction.**
 - » CPU switches to kernel mode & executes a kernel function.
- **Linux files relating to system call:**
 - » **`arch/i386/kernel/entry.S`**
 - ◆ System call and low-level fault handling routines.
 - ◆ `ENTRY(sys_call_table)`
 - » **`include/asm-i386/unistd.h`**
 - ◆ System call numbers and macros.
 - » **`kernel/sys.c`**
 - ◆ System call service routines.



Process

- **Program**

- » Structured set of commands stored in an executable file on a file system
- » Executed to create a process

- **Process**

- » Program running in memory and on the CPU (**active program**)
- » Modern systems allow 'multiprogramming' (i.e., several processes exist concurrently)
- » Each process requires an allocation of cpu time, in order to make forward progress
- » The OS must do 'scheduling' of cpu time
- » Each process also needs the use of other system facilities (memory, files, devices)
- » the terms *job* (**batch system**) and *process* is almost used interchangeably



Process

- **A process includes:**

- » text section : program
- » stack
- » data section
- » heap

process in memory

- **User process**


- » Process begun by a user that runs on a terminal (tty)

- **Daemon process**

- » System process
- » Not associated with a terminal

- **Process ID (PID) : top**

- » Unique identifier assigned to every process as it begins
- » Processes are identified by their process identifier, an integer





Process

- **Child processes**
 - » Refers to a process that was started by another process (parent process)
- **Parent processes**
 - » Process that has started other processes (child processes)
- **Parent Process ID (PPID)**
 - » The PID of the parent process that created the current process
- **Processes communicate via pipes; queues of bytes between two processes that are accessed by a file descriptor (IPC)**



Process

- **Process System Call**

- » **Creation** : fork system call creates new process
 - ◆ the process that calls fork() is parent, the new process is child
 - ◆ fork() is actually implemented via the clone() system call
 - ◆ “cloning flags” : man clone or include/linux/sched.h
- » **Execution** : exec system call used after a fork to replace the process' memory space with a new program
- » **exit** : exit terminates process and free all resources
- » **wait** : A parent may wait for a child process to terminate

- **Task**

- » Another name for a process is a task.
- » Generally refer to a process from the kernel's point of view.
- » Linux kernel internally refers to processes as tasks.



Task

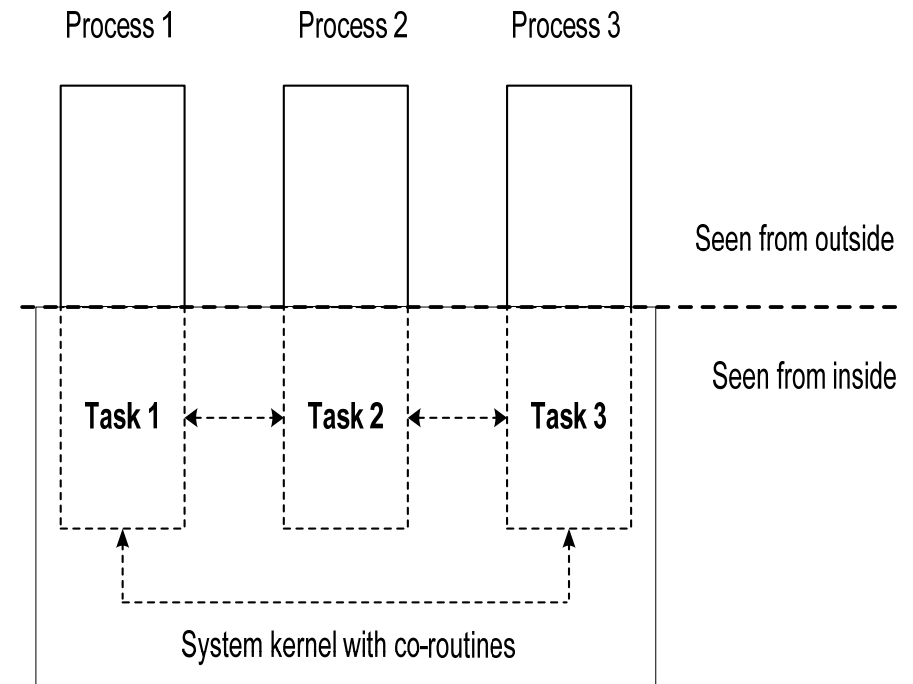
» As seen by a process running under Linux

- ◆ the kernel is a provider of services.
- ◆ Individual processes exist independently alongside each other and cannot affect each other directly.
- ◆ Each process's own area of memory is protected against modification by other processes

» The internal viewpoint of a running Linux

- ◆ Only one program (OS) is running on the computer, can access all the resource.
- ◆ The various tasks are carried out by co-routines
- ◆ an error in the kernel programming can block the entire system.

Process and Tasks





Process Control Block (PCB)

- Information associated with each process
- type sturcture **task_struct** ; line 382
» /usr/src/linux-2.4/include/linux/sched.h

Process state

Unique process identifier

Memory-management information

File system information

Signal information

Program counter

CPU registers

CPU scheduling information (e.g., priority)

Accounting information

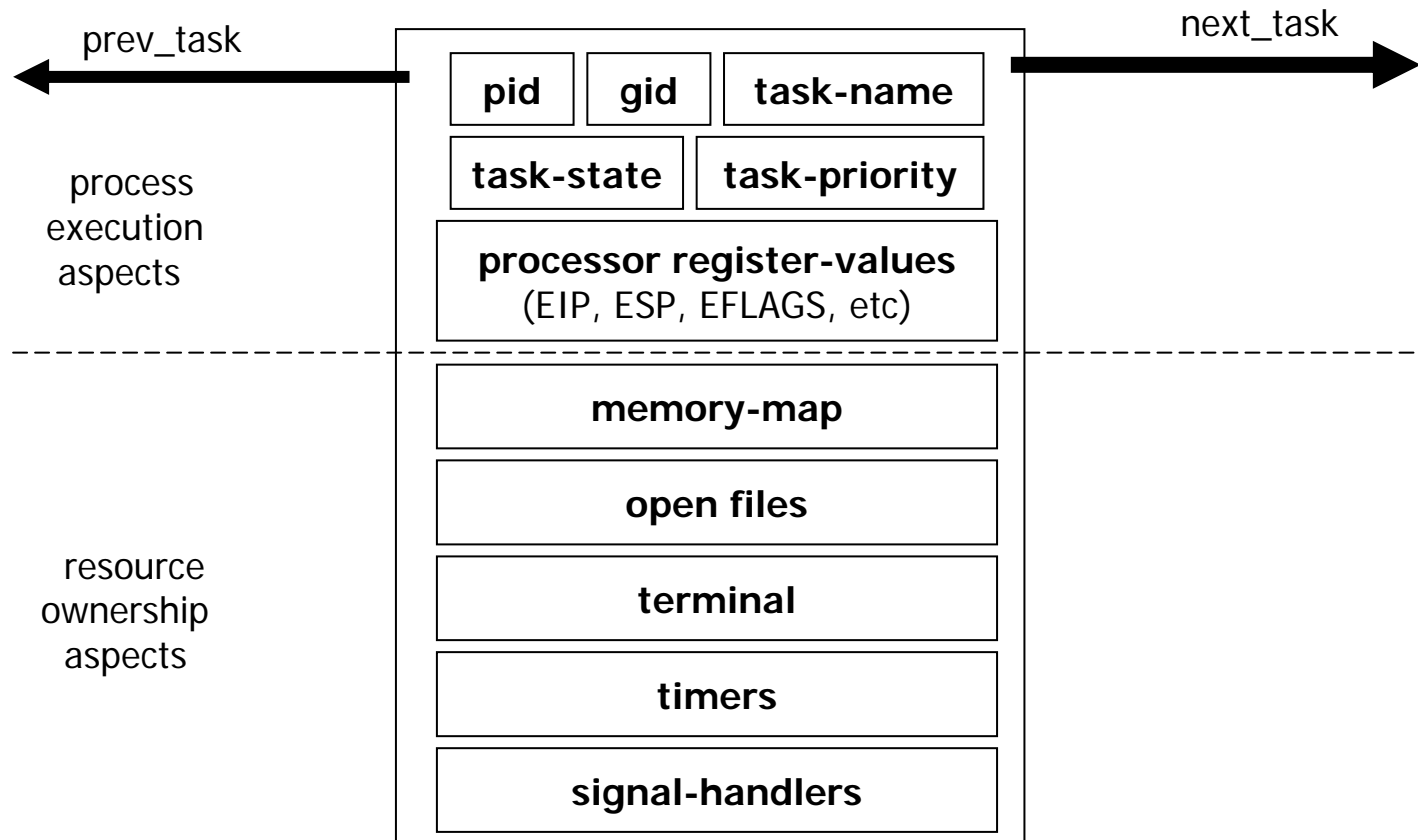
I/O status information

Pointers to other control blocks



Process Control Block (PCB)

- task_struct






Process Control Block

- task_struct

```
struct task_struct {
    volatile long state; /* -1 unrunnable, 0 runnable, >0 stopped */
    unsigned long flags; /* per process flags */
    mm_segment_t addr_limit; /* thread address space:
        0-0xBFFFFFFF for user-thread
        0-0xFFFFFFFF for kernel-thread */
    struct exec_domain *exec_domain;
    long need_resched;
    long counter;
    long priority;
    /* SMP and runqueue state */
    struct task_struct *next_task, *prev_task;
    struct task_struct *next_run, *prev_run;
    ...
    /* task state */
    /* limits */
    /* file system info */
    /* ipc stuff */
    /* tss for this task */
    ...
    /* open file information */
    /* memory management info */
    /* signal handlers */
    ...
};
```





Thread


- Thread of execution = Thread
- The object of activity within the process.
- Each thread includes a unique program counter, process stack, and set of processor registers.
- The kernel schedules individual threads, not process.
- Linux has a unique implementation of threads, and does not differentiate between threads and processes : a thread is just a special kind of process.
- A process has two distinct aspects
 - » Its 'execution' (forward progression of states)
 - » Its 'ownership' (share of system's resources)
 - » The word '**thread**' refers to the **execution aspect** of a process (i.e., the entity which gets 'scheduled' by an Operating System for a share of the available cpu time)




Thread

- **multi-threading**

- » It is possible for an operating system to support a concept of 'process' in which more than one execution-thread exists (with process-resources being shared)



one process with
one thread



one process with
multiple threads


- **Advantages of 'threads'**

- » All the 'threads' that belong to a process can access the same memory, the same files, the same terminal, the same timers and the same signal-handlers
- » Thus the system 'overhead' involved in managing a multithreaded task is more efficient (less time-consuming to set up and keep track of) than if each of those threads had individualized ownerships



Thread

- **Thread communication**
 - » Since the threads in a process share the same memory, no special mechanism is needed for them to communicate data



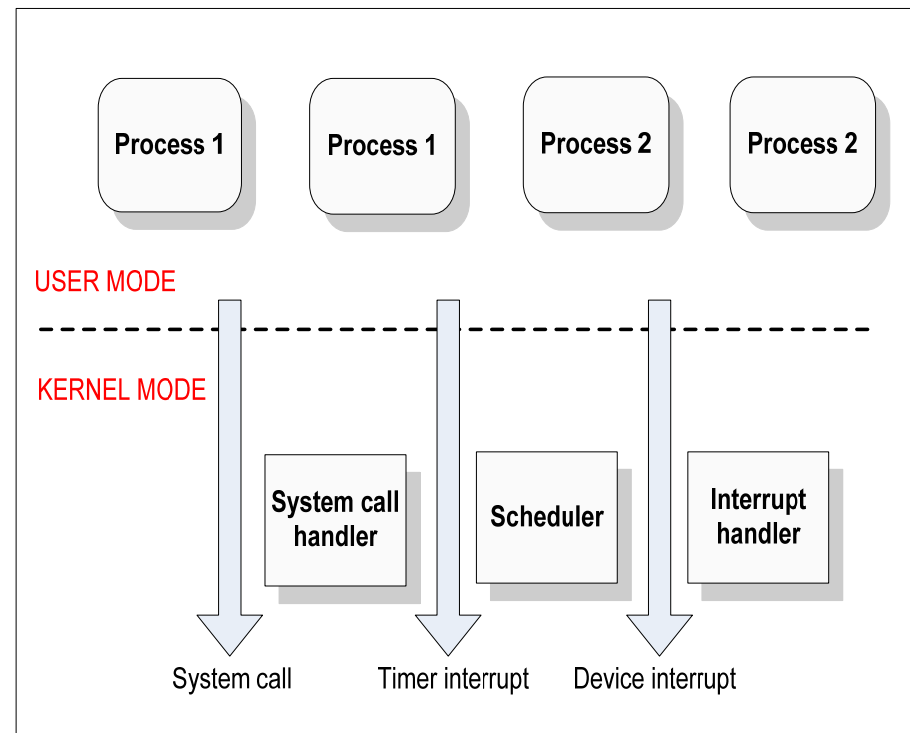
All these threads can read or write the same memory-locations



Process

- » 1. P1 in User Mode issues a system call
- » 2. the process switches to Kernel Mode, and the system call is serviced.
- » 3. P1 then resumes execution in User Mode until a timer interrupt occurs.
- » 4. the scheduler is activated in Kernel Mode.
- » 5. P2 starts its execution in User Mode until a hardware device raises an interrupt.
- » 6. As a consequence of the interrupt, P2 switches to Kernel Mode and services the interrupt.

Transitions between User and Kernel Mode





Process State

- **TASK_RUNNING** : 1. running / 2. ready
 - » The process is runnable
 - ◆ 1. currently running - running
 - ◆ 2. on a runqueue waiting to run - ready
- **TASK_INTERRUPTIBLE** : waiting/sleeping/blocked
 - » waiting on a condition: interrupts, or signals
 - » becomes runnable if it receives a signal
- **TASK_UNINTERRUPTIBLE**
 - » Waiting process cannot be woken by a signal
 - » cannot become runnable if it receives a signal
 - » this is used in situations where the process must wait without interruption.
- **TASK_STOPPED** : terminated (exit)
 - » stopped process - e.g., by a debugger
 - » this occurs if the process receives SIGSTOP signal

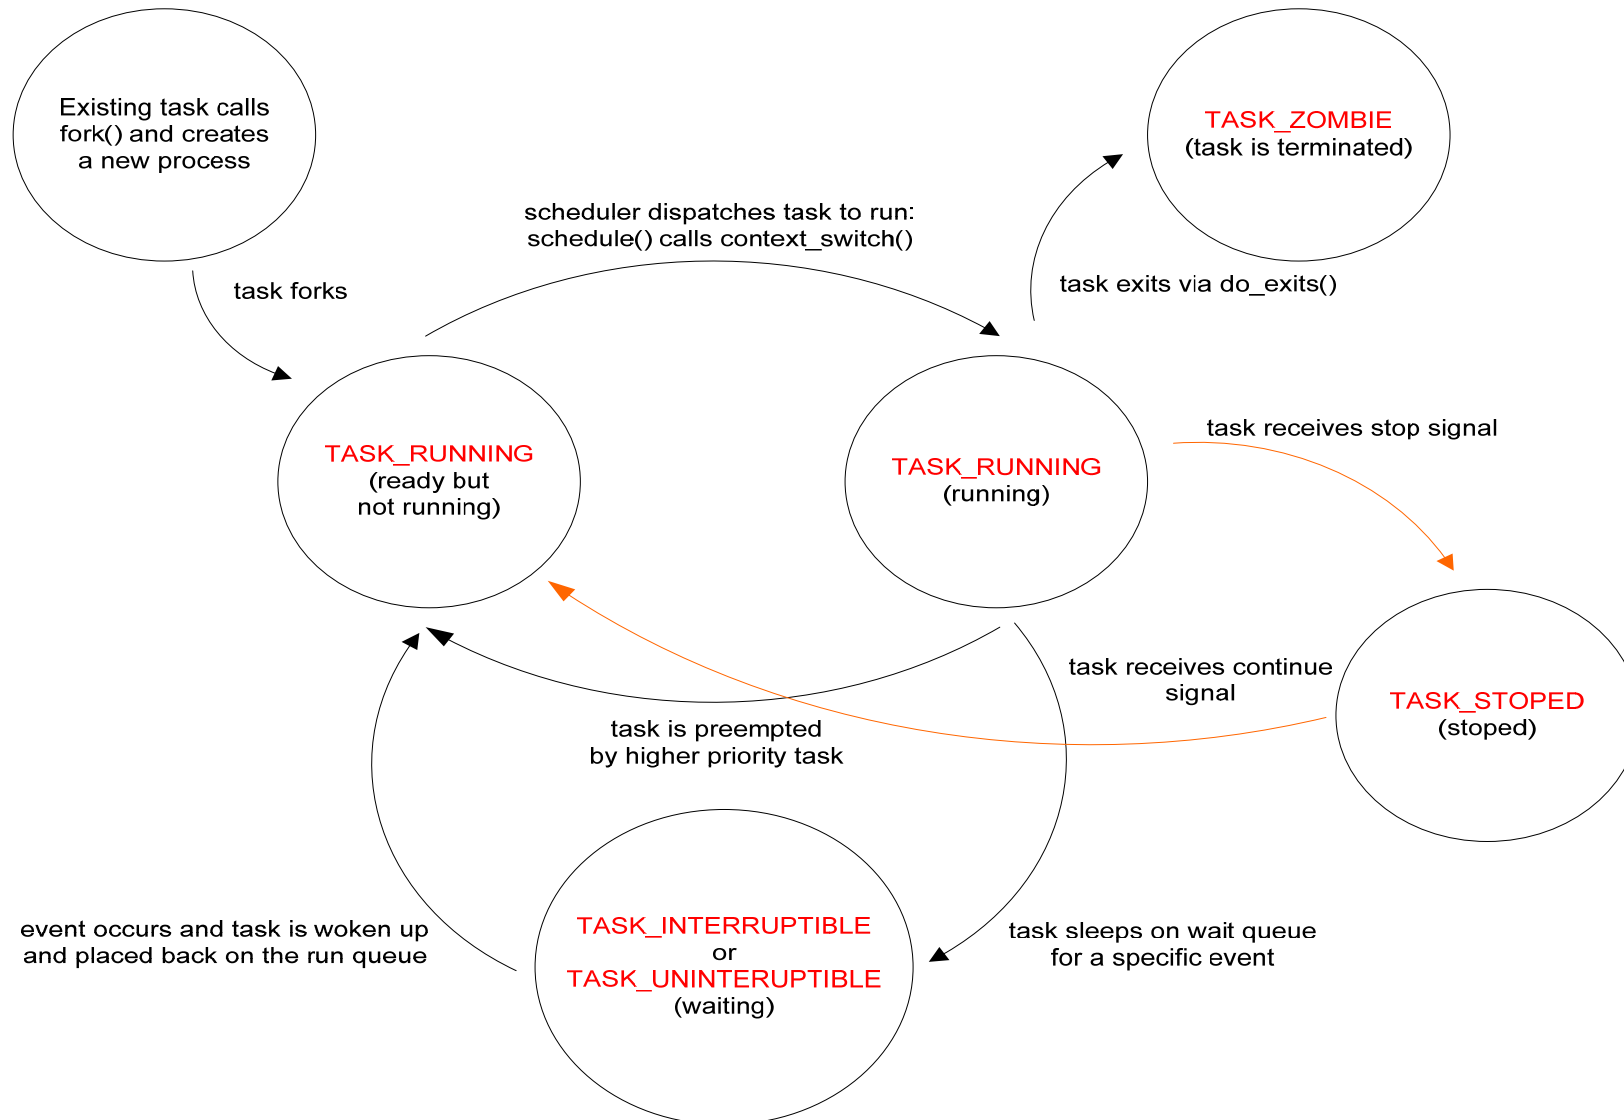


Process State

- **TASK_ZOMBIE** : defunct
 - » the process has terminated, but its parent has not yet issued a wait system call
 - » When process terminates, still consumes system resources.
 - » Process finished, but parent has not released PID.
 - » Defunct process.
 - » A parent may `wait` for a child process to terminate
 - » `wait` provides the process id of a terminated child so that the parent can tell which child terminated
 - » `wait3` allows the parent to collect performance statistics about the child
 - » Eliminated by killing the parent process



Diagram of Process State





Process

- **Context Switch**

- » When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process
- » Context-switch time is overhead; the system does no useful work while switching
- » Time dependent on hardware support

- **Remote Procedure Call**

- » Remote procedure call (RPC) abstracts procedure calls between processes on networked systems.



Process

- Viewing Processes

- » ***ps -f* command : user process**

- ◆ User identifier (UID), PID, PPID, start time, CPU utilization

- » ***ps -l* command : user process**

- ◆ More complete information
- ◆ Process state can be seen
- ◆ **R** = running, **S** = sleeping, **D** = uninterruptible sleep, **T** = stopped or being traced, **Z** = zombie

- » ***top* or *ps -efl* command : system/user process**

- ◆ Process priority (PRI) : **priority**
 - Higher value means lower priority
 - 0 (high priority) to 139 (low priority)
- ◆ Nice value (NI) : **time slice**
 - Indirectly represents priority
 - The priority value may be in the range -20 (800ms) to 19(5ms).
 - Higher value means lower priority
- ◆ **NICE_TO_PRIO(nice), PRIO_TO_NICE(prio) : kernel/sched.c**



Interprocess Communication (IPC)

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - » send(*message*) – message size fixed or variable
 - » receive(*message*)
- If *P* and *Q* wish to communicate, they need to:
 - » establish a *communication link* between them
 - » exchange messages via send/receive
- Implementation of communication link
 - » physical (e.g., shared memory, hardware bus)
 - » logical (e.g., logical properties)



Communications Model

- **Direct Communication**

- » Processes must name each other explicitly:
- » send(*P, message*) – send a message to process P
- » receive(*Q, message*) – receive a message from process Q

- **Indirect Communication**

- » Messages are directed and received from mailboxes (also referred to as ports)
- » Each mailbox has a unique id
- » Processes can communicate only if they share a mailbox
- » Operations
 - ◆ create a new mailbox
 - ◆ send and receive messages through mailbox
 - ◆ destroy a mailbox
- » send(*A, message*) – send a message to mailbox A
- » receive(*A, message*) – receive a message from mailbox A



Synchronization

- **Message passing may be either blocking or non-blocking**
- **Blocking is considered synchronous**
 - » Blocking send **has the sender block until the message is received**
 - » Blocking receive **has the receiver block until a message is available**
- **Non-blocking is considered asynchronous**
 - » Non-blocking send **has the sender send the message and continue**
 - » Non-blocking receive **has the receiver receive a valid message or null**



Signal

- Introduced in UNIX systems to simplify IPC.
- Used by the kernel to notify processes of system events.
- A signal is a short message sent to a process, or group of processes, containing the number identifying the signal.
- Linux supports 31 non-RT signals.
 - » `/include/asm-i386/signal.h`
- POSIX standard defines a range of values for RT signals



Signal

- **A signal is sent due to occurrence of corresponding event** (kernel/signal.c – line 1110)

- » `send_sig_info(int sig, struct siginfo *info, struct task_struct *t);`
- » `sig` is signal number.
- » `info` is either:
 - ♦ 0, if user mode process is signal sender.
 - ♦ 1, if kernel is signal sender.
- » `kill_proc_info(int sig, struct siginfo *info, pid_t pid);`

#	Signal Name	Default Action	Comment
1	SIGHUP	Abort	Hangup terminal or process
2	SIGINT	Abort	Keyboard interrupt (usually Ctrl-C)
...			
9	SIGKILL	Abort	Forced process termination
10	SIGUSR1	Abort	Process specific
11	SIGSEGV	Dump	Invalid memory reference
...			



Signal

- **Sending a signal**
 - » Kernel sends (delivers) a signal to a destination process by updating some state in the context of the destination process.
- **Kernel sends a signal for one of the following reasons:**
 - » **Generated internally:**
 - ◆ Kernel has detected a system event such as divide-by-zero (SIGFPE) or the termination of a child process (SIGCHLD).
 - » **Generated externally:**
 - ◆ Another process has invoked the kill system call to explicitly request the kernel to send a signal to the destination process.



Signal

- **Receiving a signal**
 - » A destination process receives a signal when it is forced by the kernel to react in some way to the delivery of the signal.
- **Three possible ways to react:**
 - » Ignore the signal (do nothing)
 - » Terminate the process
 - » Catch the signal by executing a user-level function called a signal handler.
 - ◆ Similar to a hardware exception handler being called in response to an asynchronous interrupt.



Scheduler

- **Scheduling**

- » Non-preemptive scheduling: the process keeps the CPU until the process terminates or it switches to waiting state (simple to implement).
- » Preemptive scheduling: the process can be interrupted and must release the CPU

- **Scheduling Policies** : `/linux/sched.h`

- » **SCHED_NORMAL** : non-RT / classic Unix
 - ◆ Interactive : make the system appear fast to user.
 - ◆ RT have higher priorities than any non-RT tasks.
- » **SCHED_FIFO** : soft-RT
 - ◆ if no other higher-priority RT process is runnable, the process will continue to use the CPU as long as it wishes.
- » **SCHED_RR** : soft-RT
 - ◆ assign CPU time to all SCHED_RR processes that have the same priority (fixed time slice).
 - ◆ interrupted when its time slice has expired.

* Scheduler ***schedule()*** function : `/kernel/sched.c` - line 986



Scheduler

- Long-term scheduler **(or job scheduler)**
 - » selects which processes should be brought into the ready queue
 - » invoked very infrequently (seconds, minutes) \Rightarrow (may be slow)
- Short-term scheduler **(or CPU scheduler)**
 - » selects which process should be executed next and allocates CPU
 - » invoked very frequently (milliseconds) \Rightarrow (must be fast)
- **Processes can be described as either:**
 - » I/O-bound process – spends more time doing I/O than computations, many short CPU bursts
 - » CPU-bound process – spends more time doing computations, few very long CPU bursts

Linux implicitly favors I/O-bound processes over CPU-bound ones



Module

- **Linux is a monolithic kernel**
 - » trivial modifications require kernel to be recompiled
 - » kernel is increasing in size by adding new features
 - » many modules occupy permanent space in memory though they are used rarely
- **Module: steps toward micro-kernelized Linux**
 - » small and compact kernel
 - » clean kernel
 - » rapid kernel
 - » components-based Linux
- **Easier way to extend Linux kernel functions.**
- **It is then linked into the kernel when the module is installed.**



Module

- **Most modern operating systems implement kernel modules**
 - » Uses object-oriented approach
 - » Each core component is separate
 - » Each talks to the others over known interfaces
 - » Each is loadable as needed within the kernel
- **Overall, similar to layers but with more flexible**
- **Modules can be compiled and dynamically linked into kernel address space.**
 - » Useful for device drivers that need not always be resident until needed.
 - ◆ Keeps core kernel “footprint” small.



Module versus Application

● Modules

- » 1. runs in kernel space (**no main function**).
- » 2. just registers itself in order to serve future requests, and its initialization function terminates immediately (**event-driven program**)
- » 3. no libraries to link to (linked only to the kernel): the only functions it can call are the ones exported by the kernel.
 - ◆ `printk()` : kernel library

● Applications

- » 1. runs in user space (**main function**).
- » 2. performs a single task from beginning to end (**not all applications are event-driven**)
- » 3. can call functions it doesn't define: the linking stage resolves external references using the appropriate library of functions (libc).
 - ◆ `printf()`



Module

- Every module consist of two basic functions (minimum)

```
int init_module(void) /*used for all initialization*/  
{  
    ...  
}  
void cleanup_module(void) /*used for exit */  
{  
    ...  
}
```

- Loading a module - by issuing the following command:
 » *insmod module.o*



Module

- **Instruction**

- » **insmod** : **/proc/modules**
 - ◆ Insert an module into the kernel.
- » **rmmod**
 - ◆ Remove an module from the kernel.
- » **depmod** : **/lib/modules/2.4.20-8/modules.dep**
 - ◆ Determine interdependencies between modules.
- » **ksyms** : **/proc/ksyms**
 - ◆ Display symbols that are exported by the kernel
- » **lsmod**
 - ◆ List currently loaded modules.
- » **modinfo**
 - ◆ Display module information (.modinfo section in module object file).
- » **modprobe** : **/etc/modules.conf**
 - ◆ Insert module and resolve module dependency (as described by modules.dep). For example, if you must load A before loading B, modprobe will automatically load A when you tell it to load B.



Module

- **modinfo hello.o**
 - » `MODULE_AUTHOR("MELEE");`
 - » `MODULE_DESCRIPTION("Hello Module");`
 - » `MODULE_LICENCE("GPL");`
- **/proc/ksyms : Kernel Symbol Table**
 - » lists all the symbols the kernel exports
 - » `/usr/src/linux/System.map` (core symbols)
 - » `/proc/ksyms` (all symbols)
- **/proc/modules**
 - » holds information about the loaded modules
 - » column : **module name / size / use count / dependency**
- **/etc/modules.conf**
 - » `modprobe` command
 - » A module name : `eth0`, `usb-controller`
 - » A more generic identifier : `e100`, `ehci-hcd`



Module

- **kernel 2.4 : host**

/ HELLO MODULE example -Tested on Linux Kernel 2.4 */*

```
#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/init.h>
```

```
MODULE_LICENSE("GPL");
MODULE_AUTHOR("MELEE");
MODULE_DESCRIPTION("HELLO MODULE");
```

```
int init_module(void)
{
    printk("HELLO MODULE is loaded. \n");

    return 0;
}

void cleanup_module(void)
{
    printk("HELLO MODULE is unloaded. \n");
}
```

Makefile for Linux Kernel Module for 2.4

CC = gcc

KERNEL_PATH = /usr/src/linux-2.4.20-8

CFLAGS = -DMODULE -D__KERNEL__ -I\$(KERNEL_PATH)/include

MOD_OBJ = hello

all: \$(MOD_OBJ).o

\$(MOD_OBJ): \$(MOD_OBJ).c
\$(CC) \$(CFLAGS) -c \$(MOD_OBJ).c

clean:

rm -f *.o

-D__KERNEL__ = #define __KERNEL__

-DMODULE = #define MODULE



Module

- **kernel 2.6 : target**

/ HELLO MODULE example - Tested on Linux Kernel 2.6.x */*

```
#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/init.h>
```

```
MODULE_LICENSE("GPL");
```

```
int hello_init(void)
{
    printk(" HELLO MODULE is loaded. \n");

    return 0;
}
```

```
void hello_exit(void)
{
    printk(" HELLO MODULE is unloaded. \n");
}
```

```
module_init(hello_init);
module_exit(hello_exit);
```

```
##          Makefile for Linux Kernel Module for 2.6

obj-m       := hello.o

CC          := /opt/iwmmxt-1.0.0/bin/arm-linux-gcc

KDIR        := /pxa270/kernel/linux-2.6.11-h270-tku_v1.1

PWD         := $(shell pwd)

default:
    $(MAKE) -C $(KDIR) SUBDIRS=$(PWD) modules

clean:
    rm -rf *.ko
    rm -rf *.mod.*
    rm -rf *.cmd
    rm -rf .tmp*
    rm -rf *.o
```




Module

- Make module : **sample modules** (microcom homepage)
- Load the module
 - » *insmod hello.o* or *hello.ko*
 - » *dmesg* : at host - /var/log/messages
- Check that the module is loaded
 - » *cat /proc/modules*
- Remove the module
 - » *rmmmod hello*